

Download the necessary libraries

```
In [500... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sn
```

Load the Dataset

```
In [501... df = pd.read_csv("C:\\Users\\captr\\OneDrive\\Desktop\\AusAppara1Sales4thQrt2020.csv")
df.head(5)
```

```
Out[501]:
```

	Date	Time	State	Group	Unit	Sales
0	1-Oct-2020	Morning	WA	Kids	8	20000
1	1-Oct-2020	Morning	WA	Men	8	20000
2	1-Oct-2020	Morning	WA	Women	4	10000
3	1-Oct-2020	Morning	WA	Seniors	15	37500
4	1-Oct-2020	Afternoon	WA	Kids	3	7500

```
In [502... type(df['Date'])
```

```
Out[502]: pandas.core.series.Series
```

Checking for the missing values

```
In [503... df.isnull().sum()
```

```
Out[503]: Date      0
Time      0
State     0
Group     0
Unit      0
Sales     0
dtype: int64
```

Checking for the duplicate values (if any)

```
In [504... df.duplicated().sum()
```

```
Out[504]: 0
```

Running Basic Summary Statistics of the data

```
In [505... df.describe().T
```

Out[505]:

	count	mean	std	min	25%	50%	75%	max
Unit	7560.0	18.005423	12.901403	2.0	8.0	14.0	26.0	65.0
Sales	7560.0	45013.558201	32253.506944	5000.0	20000.0	35000.0	65000.0	162500.0

Checking how many times different states of Australia appeared in our dataset

In [506...]

```
state_counts=df['State'].value_counts()
state_counts
```

Out[506]:

```
State
WA      1080
NT      1080
SA      1080
VIC      1080
QLD      1080
NSW      1080
TAS      1080
Name: count, dtype: int64
```

In [507...]

```
state_counts.sort_values()
```

Out[507]:

```
State
WA      1080
NT      1080
SA      1080
VIC      1080
QLD      1080
NSW      1080
TAS      1080
Name: count, dtype: int64
```

Two of our features i.e Unit and Sales were quite wide apart in terms of their range hence performed min-max-scaling to transform it to a fixed range between 0 to 1 or [0,1]

In [508...]

```
from sklearn.preprocessing import MinMaxScaler
```

In [509...]


```
scaler= MinMaxScaler()
scaler.fit(df[['Sales']])
df['Sales']= scaler.transform(df[['Sales']])
scaler.fit(df[['Unit']])
df['Unit']= scaler.transform(df[['Unit']])
df.head()
```

Out[509]:

	Date	Time	State	Group	Unit	Sales
0	1-Oct-2020	Morning	WA	Kids	0.095238	0.095238
1	1-Oct-2020	Morning	WA	Men	0.095238	0.095238
2	1-Oct-2020	Morning	WA	Women	0.031746	0.031746
3	1-Oct-2020	Morning	WA	Seniors	0.206349	0.206349
4	1-Oct-2020	Afternoon	WA	Kids	0.015873	0.015873

Plotted Unit against time

```
In [510... %matplotlib inline
plt.scatter(df['Unit'],df['Sales'], color='red', marker='+')
plt.xlabel('Units')
plt.ylabel('Sales Quantity')
plt.show()
```



Grouped the data according to the State/region

```
In [511... g=df.groupby('State')
g.head(5)
```

Out[511]:

	Date	Time	State	Group	Unit	Sales
0	1-Oct-2020	Morning	WA	Kids	0.095238	0.095238
1	1-Oct-2020	Morning	WA	Men	0.095238	0.095238
2	1-Oct-2020	Morning	WA	Women	0.031746	0.031746
3	1-Oct-2020	Morning	WA	Seniors	0.206349	0.206349
4	1-Oct-2020	Afternoon	WA	Kids	0.015873	0.015873
12	1-Oct-2020	Morning	NT	Kids	0.174603	0.174603
13	1-Oct-2020	Morning	NT	Men	0.047619	0.047619
14	1-Oct-2020	Morning	NT	Women	0.031746	0.031746
15	1-Oct-2020	Morning	NT	Seniors	0.126984	0.126984
16	1-Oct-2020	Afternoon	NT	Kids	0.174603	0.174603
24	1-Oct-2020	Morning	SA	Kids	0.158730	0.158730
25	1-Oct-2020	Morning	SA	Men	0.158730	0.158730
26	1-Oct-2020	Morning	SA	Women	0.222222	0.222222
27	1-Oct-2020	Morning	SA	Seniors	0.412698	0.412698
28	1-Oct-2020	Afternoon	SA	Kids	0.222222	0.222222
36	1-Oct-2020	Morning	VIC	Kids	0.746032	0.746032
37	1-Oct-2020	Morning	VIC	Men	0.539683	0.539683
38	1-Oct-2020	Morning	VIC	Women	0.507937	0.507937
39	1-Oct-2020	Morning	VIC	Seniors	0.380952	0.380952
40	1-Oct-2020	Afternoon	VIC	Kids	0.460317	0.460317
48	1-Oct-2020	Morning	QLD	Kids	0.285714	0.285714
49	1-Oct-2020	Morning	QLD	Men	0.253968	0.253968
50	1-Oct-2020	Morning	QLD	Women	0.206349	0.206349
51	1-Oct-2020	Morning	QLD	Seniors	0.190476	0.190476
52	1-Oct-2020	Afternoon	QLD	Kids	0.253968	0.253968
60	1-Oct-2020	Morning	NSW	Kids	0.587302	0.587302
61	1-Oct-2020	Morning	NSW	Men	0.238095	0.238095
62	1-Oct-2020	Morning	NSW	Women	0.507937	0.507937
63	1-Oct-2020	Morning	NSW	Seniors	0.333333	0.333333
64	1-Oct-2020	Afternoon	NSW	Kids	0.603175	0.603175
72	1-Oct-2020	Morning	TAS	Kids	0.174603	0.174603
73	1-Oct-2020	Morning	TAS	Men	0.063492	0.063492
74	1-Oct-2020	Morning	TAS	Women	0.015873	0.015873
75	1-Oct-2020	Morning	TAS	Seniors	0.095238	0.095238
76	1-Oct-2020	Afternoon	TAS	Kids	0.079365	0.079365

In [512...

```
for i, i_df in g:  
    print(i)  
    print(i_df)
```

NSW						
	Date	Time	State	Group	Unit	Sales
60	1-Oct-2020	Morning	NSW	Kids	0.587302	0.587302
61	1-Oct-2020	Morning	NSW	Men	0.238095	0.238095
62	1-Oct-2020	Morning	NSW	Women	0.507937	0.507937
63	1-Oct-2020	Morning	NSW	Seniors	0.333333	0.333333
64	1-Oct-2020	Afternoon	NSW	Kids	0.603175	0.603175
...
7543	30-Dec-2020	Afternoon	NSW	Seniors	0.269841	0.269841
7544	30-Dec-2020	Evening	NSW	Kids	0.555556	0.555556
7545	30-Dec-2020	Evening	NSW	Men	0.619048	0.619048
7546	30-Dec-2020	Evening	NSW	Women	0.555556	0.555556
7547	30-Dec-2020	Evening	NSW	Seniors	0.333333	0.333333

[1080 rows x 6 columns]

NT						
	Date	Time	State	Group	Unit	Sales
12	1-Oct-2020	Morning	NT	Kids	0.174603	0.174603
13	1-Oct-2020	Morning	NT	Men	0.047619	0.047619
14	1-Oct-2020	Morning	NT	Women	0.031746	0.031746
15	1-Oct-2020	Morning	NT	Seniors	0.126984	0.126984
16	1-Oct-2020	Afternoon	NT	Kids	0.174603	0.174603
...
7495	30-Dec-2020	Afternoon	NT	Seniors	0.174603	0.174603
7496	30-Dec-2020	Evening	NT	Kids	0.206349	0.206349
7497	30-Dec-2020	Evening	NT	Men	0.063492	0.063492
7498	30-Dec-2020	Evening	NT	Women	0.142857	0.142857
7499	30-Dec-2020	Evening	NT	Seniors	0.190476	0.190476

[1080 rows x 6 columns]

QLD						
	Date	Time	State	Group	Unit	Sales
48	1-Oct-2020	Morning	QLD	Kids	0.285714	0.285714
49	1-Oct-2020	Morning	QLD	Men	0.253968	0.253968
50	1-Oct-2020	Morning	QLD	Women	0.206349	0.206349
51	1-Oct-2020	Morning	QLD	Seniors	0.190476	0.190476
52	1-Oct-2020	Afternoon	QLD	Kids	0.253968	0.253968
...
7531	30-Dec-2020	Afternoon	QLD	Seniors	0.111111	0.111111
7532	30-Dec-2020	Evening	QLD	Kids	0.174603	0.174603
7533	30-Dec-2020	Evening	QLD	Men	0.365079	0.365079
7534	30-Dec-2020	Evening	QLD	Women	0.285714	0.285714
7535	30-Dec-2020	Evening	QLD	Seniors	0.253968	0.253968

[1080 rows x 6 columns]

SA						
	Date	Time	State	Group	Unit	Sales
24	1-Oct-2020	Morning	SA	Kids	0.158730	0.158730
25	1-Oct-2020	Morning	SA	Men	0.158730	0.158730
26	1-Oct-2020	Morning	SA	Women	0.222222	0.222222
27	1-Oct-2020	Morning	SA	Seniors	0.412698	0.412698
28	1-Oct-2020	Afternoon	SA	Kids	0.222222	0.222222
...
7507	30-Dec-2020	Afternoon	SA	Seniors	0.269841	0.269841
7508	30-Dec-2020	Evening	SA	Kids	0.460317	0.460317
7509	30-Dec-2020	Evening	SA	Men	0.206349	0.206349
7510	30-Dec-2020	Evening	SA	Women	0.507937	0.507937
7511	30-Dec-2020	Evening	SA	Seniors	0.507937	0.507937

[1080 rows x 6 columns]

TAS						
	Date	Time	State	Group	Unit	Sales
72	1-Oct-2020	Morning	TAS	Kids	0.174603	0.174603
73	1-Oct-2020	Morning	TAS	Men	0.063492	0.063492

74	1-Oct-2020	Morning	TAS	Women	0.015873	0.015873
75	1-Oct-2020	Morning	TAS	Seniors	0.095238	0.095238
76	1-Oct-2020	Afternoon	TAS	Kids	0.079365	0.079365
...
7555	30-Dec-2020	Afternoon	TAS	Seniors	0.190476	0.190476
7556	30-Dec-2020	Evening	TAS	Kids	0.206349	0.206349
7557	30-Dec-2020	Evening	TAS	Men	0.206349	0.206349
7558	30-Dec-2020	Evening	TAS	Women	0.142857	0.142857
7559	30-Dec-2020	Evening	TAS	Seniors	0.174603	0.174603

[1080 rows x 6 columns]

VIC

	Date	Time	State	Group	Unit	Sales
36	1-Oct-2020	Morning	VIC	Kids	0.746032	0.746032
37	1-Oct-2020	Morning	VIC	Men	0.539683	0.539683
38	1-Oct-2020	Morning	VIC	Women	0.507937	0.507937
39	1-Oct-2020	Morning	VIC	Seniors	0.380952	0.380952
40	1-Oct-2020	Afternoon	VIC	Kids	0.460317	0.460317
...
7519	30-Dec-2020	Afternoon	VIC	Seniors	0.952381	0.952381
7520	30-Dec-2020	Evening	VIC	Kids	0.444444	0.444444
7521	30-Dec-2020	Evening	VIC	Men	0.523810	0.523810
7522	30-Dec-2020	Evening	VIC	Women	0.444444	0.444444
7523	30-Dec-2020	Evening	VIC	Seniors	0.476190	0.476190

[1080 rows x 6 columns]

WA

	Date	Time	State	Group	Unit	Sales
0	1-Oct-2020	Morning	WA	Kids	0.095238	0.095238
1	1-Oct-2020	Morning	WA	Men	0.095238	0.095238
2	1-Oct-2020	Morning	WA	Women	0.031746	0.031746
3	1-Oct-2020	Morning	WA	Seniors	0.206349	0.206349
4	1-Oct-2020	Afternoon	WA	Kids	0.015873	0.015873
...
7483	30-Dec-2020	Afternoon	WA	Seniors	0.206349	0.206349
7484	30-Dec-2020	Evening	WA	Kids	0.063492	0.063492
7485	30-Dec-2020	Evening	WA	Men	0.063492	0.063492
7486	30-Dec-2020	Evening	WA	Women	0.190476	0.190476
7487	30-Dec-2020	Evening	WA	Seniors	0.111111	0.111111

[1080 rows x 6 columns]

Performed one hot encoding to convert the categorical variables in Unit and Sales column to machine readable numeric form

```
In [513... from sklearn.preprocessing import LabelEncoder
le_Time=LabelEncoder()
le_State= LabelEncoder()
le_Group = LabelEncoder()
```

```
In [514... df['Time_n']= le_Time.fit_transform(df['Time'])
df['State_n']= le_State.fit_transform(df['State'])
df['Group_n']= le_Group.fit_transform(df['Group'])
df.head()
```

Out[514]:

	Date	Time	State	Group	Unit	Sales	Time_n	State_n	Group_n
0	1-Oct-2020	Morning	WA	Kids	0.095238	0.095238	2	6	0
1	1-Oct-2020	Morning	WA	Men	0.095238	0.095238	2	6	1
2	1-Oct-2020	Morning	WA	Women	0.031746	0.031746	2	6	3
3	1-Oct-2020	Morning	WA	Seniors	0.206349	0.206349	2	6	2
4	1-Oct-2020	Afternoon	WA	Kids	0.015873	0.015873	0	6	0

In [515...]

```
df.head()
df = df.reset_index() # Resets the index, turning 'Date' back into a column
print(df.head()) # Check the first few rows
```

	index	Date	Time	State	Group	Unit	Sales	Time_n	\
0	0	1-Oct-2020	Morning	WA	Kids	0.095238	0.095238	2	
1	1	1-Oct-2020	Morning	WA	Men	0.095238	0.095238	2	
2	2	1-Oct-2020	Morning	WA	Women	0.031746	0.031746	2	
3	3	1-Oct-2020	Morning	WA	Seniors	0.206349	0.206349	2	
4	4	1-Oct-2020	Afternoon	WA	Kids	0.015873	0.015873	0	

	State_n	Group_n
0	6	0
1	6	1
2	6	3
3	6	2
4	6	0

Calculated the mean and Standard Deviation of the Unit Column

In [516...]

```
mean_value = df['Unit'].mean()
median_value = df['Unit'].median()
std_value = df['Unit'].std()

print("Mean of Unit column:", mean_value)
print("Median of Unit column:", median_value)
print("Standard Deviation of Unit column:", std_value)
```

Mean of Unit column: 0.2540543377844965
 Median of Unit column: 0.19047619047619047
 Standard Deviation of Unit column: 0.20478417107280086

Calculated the mean and standard deviation of the Sales Column

In [517...]

```
mean_value = df['Sales'].mean()
median_value = df['Sales'].median()
std_value = df['Sales'].std()

print("Mean of Sales column:", mean_value)
print("Median of Sales column:", median_value)
print("Standard Deviation of Sales column:", std_value)
```

Mean of Sales column: 0.25405433778449654
 Median of Sales column: 0.1904761904761905
 Standard Deviation of Sales column: 0.2047841710728009

Identifying the group with the highest sales and lowest sales respectively

```
In [518... # Group by 'Group' and calculate total sales
group_sales = df.groupby('Group_n')['Sales'].sum()

# Identify the group with the highest and lowest sales
highest_sales_group = group_sales.idxmax()
lowest_sales_group = group_sales.idxmin()

# Get the sales values
highest_sales_value = group_sales.max()
lowest_sales_value = group_sales.min()

print(f"Group with the highest sales: {highest_sales_group} ({highest_sales_value})")
print(f"Group with the lowest sales: {lowest_sales_group} ({lowest_sales_value})")

Group with the highest sales: 1 (484.44444444444446)
Group with the lowest sales: 2 (473.57142857142856)
```

```
In [519... df.columns

Out[519]: Index(['index', 'Date', 'Time', 'State', 'Group', 'Unit', 'Sales', 'Time_n',
      'State_n', 'Group_n'],
      dtype='object')
```

Gathering the weekly, Monthly and quarterly report

```
In [520... import pandas as pd

# Assuming `df` contains 'Date' and 'Sales' columns
df['Date'] = pd.to_datetime(df['Date']) # Ensure the Date column is in datetime format

# Set the Date column as the index for resampling
df.set_index('Date', inplace=True)

# Weekly Report
weekly_report = df.resample('W').sum()

# Monthly Report
monthly_report = df.resample('M').sum()

# Quarterly Report
quarterly_report = df.resample('Q').sum()

# Reset index for reports (optional, for better formatting)
weekly_report.reset_index(inplace=True)
monthly_report.reset_index(inplace=True)
quarterly_report.reset_index(inplace=True)

# Save reports to files
weekly_report.to_csv('weekly_report.csv', index=False)
monthly_report.to_csv('monthly_report.csv', index=False)
quarterly_report.to_csv('quarterly_report.csv', index=False)

# Print summaries
print("Weekly Report:\n", weekly_report.head())
print("Monthly Report:\n", monthly_report.head())
print("Quarterly Report:\n", quarterly_report.head())
```

Weekly Report:

	Date	index					Time	\
0	2020-10-04	56280	Morning	Morning	Morning	Morning	Afternoon	Aft...
1	2020-10-11	370146	Morning	Morning	Morning	Morning	Afternoon	Aft...
2	2020-10-18	715890	Morning	Morning	Morning	Morning	Afternoon	Aft...
3	2020-10-25	1061634	Morning	Morning	Morning	Morning	Afternoon	Aft...
4	2020-11-01	1185156	Morning	Morning	Morning	Morning	Afternoon	Aft...

	State												\			
0	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	NT	NT	NT	...
1	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	NT	NT	NT	...
2	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	NT	NT	NT	...
3	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	NT	NT	NT	...
4	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	NT	NT	NT	...

					Group					Unit			Sales	\
0	Kids	Men	Women	Seniors	Kids	Men	Women	Seniors...	84.857143			84.857143		
1	Kids	Men	Women	Seniors	Kids	Men	Women	Seniors...	152.777778			152.777778		
2	Kids	Men	Women	Seniors	Kids	Men	Women	Seniors...	150.476190			150.476190		
3	Kids	Men	Women	Seniors	Kids	Men	Women	Seniors...	151.587302			151.587302		
4	Kids	Men	Women	Seniors	Kids	Men	Women	Seniors...	122.460317			122.460317		

	Time_n	State_n	Group_n
0	336	1008	504
1	588	1764	882
2	588	1764	882
3	588	1764	882
4	504	1512	756

Monthly Report:

	Date	index					Time	\
0	2020-10-31	3173940	Morning	Morning	Morning	Morning	Afternoon	Aft...
1	2020-11-30	9524340	Morning	Morning	Morning	Morning	Afternoon	Aft...
2	2020-12-31	15874740	Morning	Morning	Morning	Morning	Afternoon	Aft...

	State												\			
0	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	NT	NT	NT	...
1	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	NT	NT	NT	...
2	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	NT	NT	NT	...

					Group					Unit			Sales	\
0	Kids	Men	Women	Seniors	Kids	Men	Women	Seniors...	645.650794			645.650794		
1	Kids	Men	Women	Seniors	Kids	Men	Women	Seniors...	495.761905			495.761905		
2	Kids	Men	Women	Seniors	Kids	Men	Women	Seniors...	779.238095			779.238095		

	Time_n	State_n	Group_n
0	2520	7560	3780
1	2520	7560	3780
2	2520	7560	3780

Quarterly Report:

	Date	index					Time	\
0	2020-12-31	28573020	Morning	Morning	Morning	Morning	Afternoon	Aft...

	State												\			
0	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	WA	NT	NT	NT	...

					Group					Unit	\
0	Kids	Men	Women	Seniors	Kids	Men	Women	Seniors...	1920.650794		

	Sales	Time_n	State_n	Group_n
0	1920.650794	7560	22680	11340

```
In [521... df_new= pd.read_csv("C:\\Users\\captr\\OneDrive\\Desktop\\AusApparalSales4thQrt2020")
df_new.head(3)
```

Out[521]:

	Date	Time	State	Group	Unit	Sales
0	1-Oct-2020	Morning	WA	Kids	8	20000
1	1-Oct-2020	Morning	WA	Men	8	20000
2	1-Oct-2020	Morning	WA	Women	4	10000

Plotted Group Vs States

In [522...

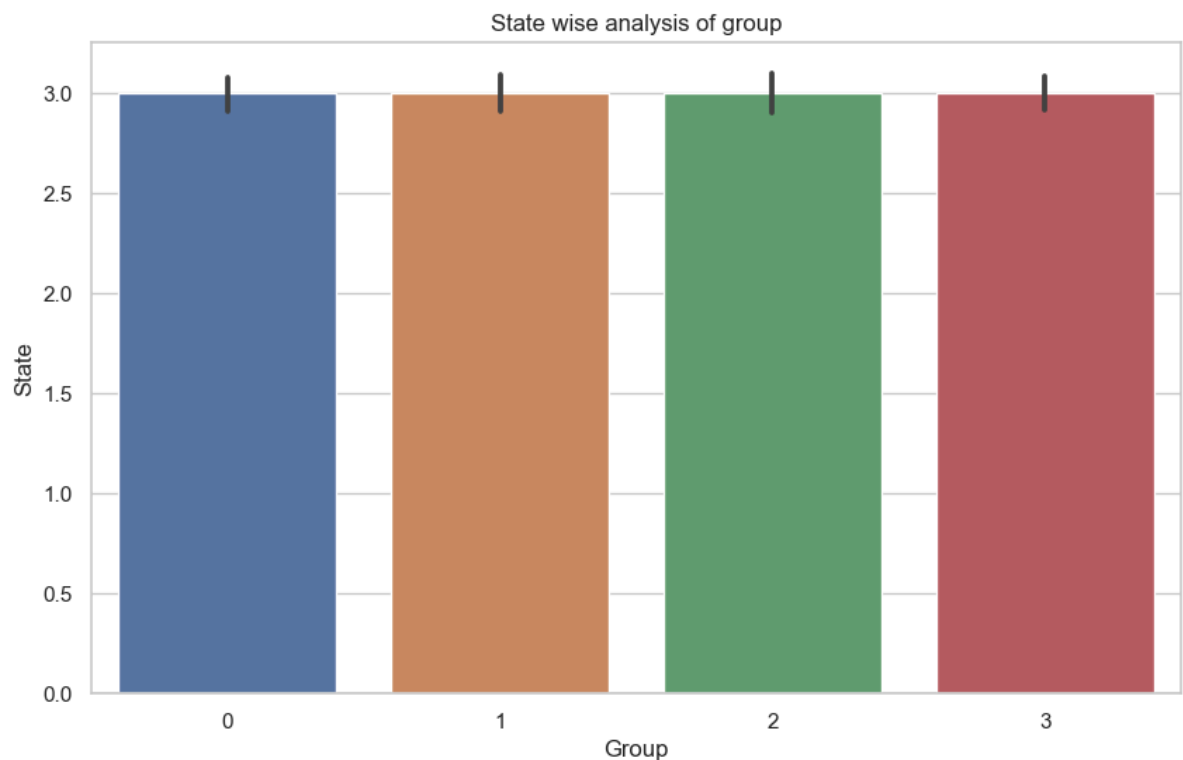
```
Group_counts= df_new['Group'].value_counts()  
Group_counts  
top_groups=Group_counts.head(10)  
top_groups
```

Out[522]:

```
Group  
Kids      1890  
Men       1890  
Women     1890  
Seniors   1890  
Name: count, dtype: int64
```

In [523...

```
import seaborn as sns  
plt.figure(figsize=(10,6))  
sns.barplot(x=df['Group_n'], y=df['State_n'])  
plt.xlabel('Group')  
plt.ylabel('State')  
plt.title('State wise analysis of group')  
plt.show()
```



In [524...

```
demography= df.groupby(['Group_n','State_n'])  
df=df.reset_index()  
df.head()
```

Out[524]:

	Date	index	Time	State	Group	Unit	Sales	Time_n	State_n	Group_n
0	2020-10-01	0	Morning	WA	Kids	0.095238	0.095238	2	6	0
1	2020-10-01	1	Morning	WA	Men	0.095238	0.095238	2	6	1
2	2020-10-01	2	Morning	WA	Women	0.031746	0.031746	2	6	3
3	2020-10-01	3	Morning	WA	Seniors	0.206349	0.206349	2	6	2
4	2020-10-01	4	Afternoon	WA	Kids	0.015873	0.015873	0	6	0

State-wise Sales Analysis for Different Demographic Groups

In [525...]

```
# Mapping Group Codes to Demographics
group_mapping = {0: 'Kids', 1: 'Women', 2: 'Men', 3: 'Seniors'}
df['Demographic'] = df['Group_n'].map(group_mapping)
group_mapping_2 = {0: 'TAS', 1: 'NSW', 2: 'QLD', 3: 'VIC', 4: 'SA', 5: 'NT', 6: 'WA'}
df['Region'] = df['State_n'].map(group_mapping_2)

sales_data = df.groupby(['State_n', 'Demographic'])['Sales'].sum().reset_index()
sales_data
```

Out[525]:

	State_n	Demographic	Sales
0	0	Kids	109.444444
1	0	Men	106.904762
2	0	Seniors	113.158730
3	0	Women	112.206349
4	1	Kids	27.619048
5	1	Men	26.126984
6	1	Seniors	27.317460
7	1	Women	28.015873
8	2	Kids	45.460317
9	2	Men	43.428571
10	2	Seniors	44.285714
11	2	Women	44.714286
12	3	Kids	83.587302
13	3	Men	84.873016
14	3	Seniors	86.476190
15	3	Women	84.476190
16	4	Kids	28.095238
17	4	Men	27.301587
18	4	Seniors	26.841270
19	4	Women	27.984127
20	5	Kids	158.793651
21	5	Men	158.507937
22	5	Seniors	159.571429
23	5	Women	159.095238
24	6	Kids	27.142857
25	6	Men	26.428571
26	6	Seniors	24.841270
27	6	Women	27.952381

In [526...]

```
import seaborn as sns
import matplotlib.pyplot as plt

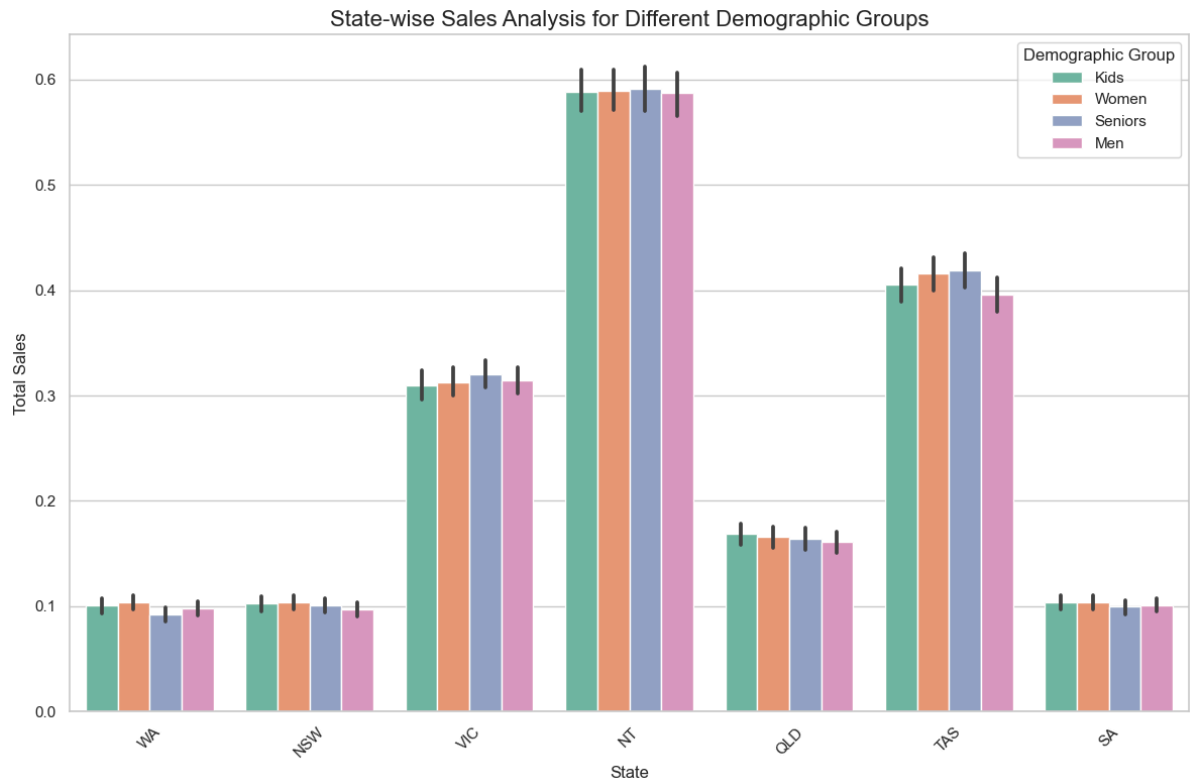
# Set plot style
sns.set(style="whitegrid")

# Create the bar plot
plt.figure(figsize=(12, 8))
sns.barplot(x='Region', y='Sales', hue='Demographic', data=df, palette="Set2")

# Customize plot
plt.title("State-wise Sales Analysis for Different Demographic Groups", fontsize=16)
```

```
plt.xlabel("State", fontsize=12)
plt.ylabel("Total Sales", fontsize=12)
plt.legend(title="Demographic Group")
plt.xticks(rotation=45)
plt.tight_layout()

# Show plot
plt.show()
```



In [527... `df['State_n'].unique()`

Out[527]: `array([6, 1, 3, 5, 2, 0, 4])`

In [528... `df.tail()`

	Date	index	Time	State	Group	Unit	Sales	Time_n	State_n	Group_n	Den
7555	2020-12-30	7555	Afternoon	TAS	Seniors	0.190476	0.190476	0	4	2	
7556	2020-12-30	7556	Evening	TAS	Kids	0.206349	0.206349	1	4	0	
7557	2020-12-30	7557	Evening	TAS	Men	0.206349	0.206349	1	4	1	
7558	2020-12-30	7558	Evening	TAS	Women	0.142857	0.142857	1	4	3	
7559	2020-12-30	7559	Evening	TAS	Seniors	0.174603	0.174603	1	4	2	

State-wise Sales Analysis for Different Regions

```
In [529... group_mapping = {0: 'Kids', 1: 'Women', 2: 'Men', 3: 'Seniors'}
df['Demographic'] = df['Group_n'].map(group_mapping)
```

```
group_mapping_2 = {0: 'TAS', 1: 'NSW', 2: 'QLD', 3: 'VIC', 4: 'SA', 5: 'NT', 6: 'WA'}
df['Region'] = df['State_n'].map(group_mapping_2)

sales_data_2 = df.groupby(['Group_n', 'Region'])['Sales'].sum().reset_index()
df.head()
```

Out[529]:

	Date	index	Time	State	Group	Unit	Sales	Time_n	State_n	Group_n	Demog
0	2020-10-01	0	Morning	WA	Kids	0.095238	0.095238	2	6	0	
1	2020-10-01	1	Morning	WA	Men	0.095238	0.095238	2	6	1	W
2	2020-10-01	2	Morning	WA	Women	0.031746	0.031746	2	6	3	S
3	2020-10-01	3	Morning	WA	Seniors	0.206349	0.206349	2	6	2	
4	2020-10-01	4	Afternoon	WA	Kids	0.015873	0.015873	0	6	0	



In [530...

```
sales_data_2
```

Out[530]:

	Group_n	Region	Sales
0	0	NSW	27.619048
1	0	NT	158.793651
2	0	QLD	45.460317
3	0	SA	28.095238
4	0	TAS	109.444444
5	0	VIC	83.587302
6	0	WA	27.142857
7	1	NSW	28.015873
8	1	NT	159.095238
9	1	QLD	44.714286
10	1	SA	27.984127
11	1	TAS	112.206349
12	1	VIC	84.476190
13	1	WA	27.952381
14	2	NSW	26.126984
15	2	NT	158.507937
16	2	QLD	43.428571
17	2	SA	27.301587
18	2	TAS	106.904762
19	2	VIC	84.873016
20	2	WA	26.428571
21	3	NSW	27.317460
22	3	NT	159.571429
23	3	QLD	44.285714
24	3	SA	26.841270
25	3	TAS	113.158730
26	3	VIC	86.476190
27	3	WA	24.841270

In [531]...

```

# Ensure 'sales_data_2' has the correct data
print(sales_data_2.head()) # Check the first few rows

# Create the bar plot
plt.figure(figsize=(12, 8))
sns.barplot(x='Demographic', y='Sales', hue='Region', data=df, palette="Set2")

# Customize the plot
plt.title("State-wise Sales Analysis for Different Regions", fontsize=16)
plt.xlabel("Group", fontsize=12)
plt.ylabel("Total Sales", fontsize=12)
plt.legend(title="Region")

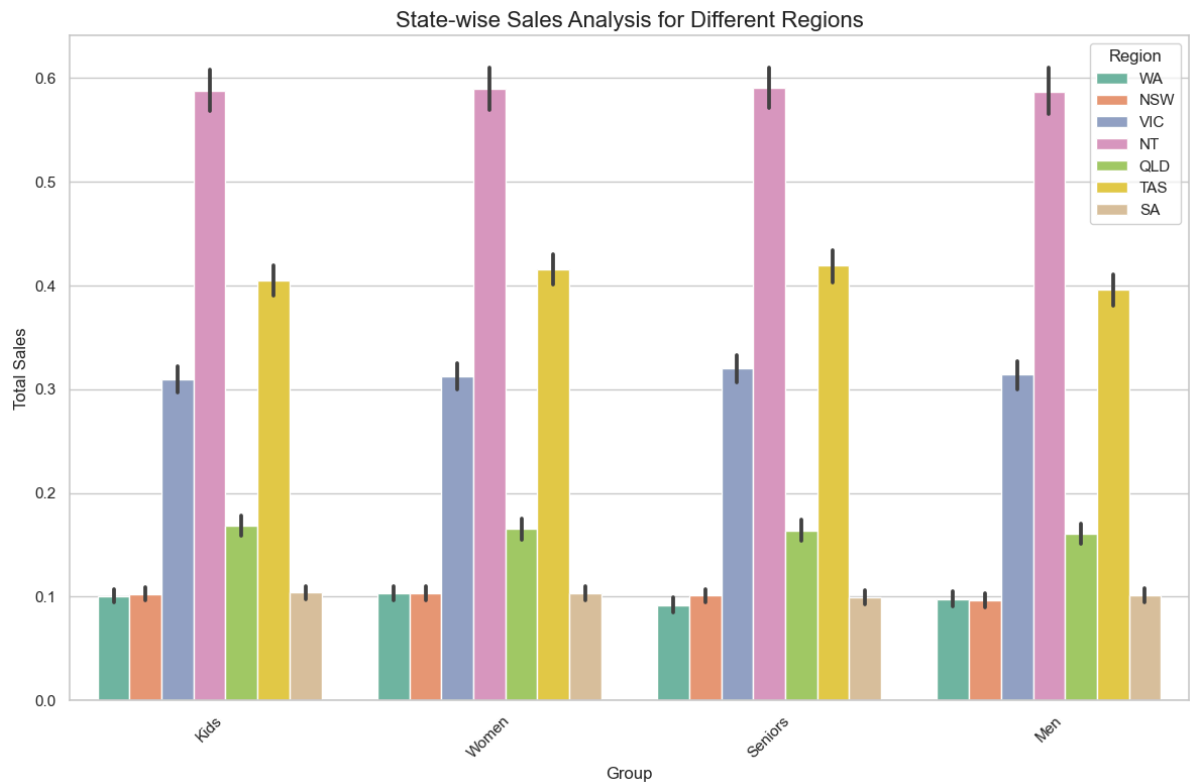
```



```
plt.xticks(rotation=45)
plt.tight_layout()
```

```
# Show plot
plt.show()
```

	Group_n	Region	Sales
0	0	NSW	27.619048
1	0	NT	158.793651
2	0	QLD	45.460317
3	0	SA	28.095238
4	0	TAS	109.444444



Sales Analysis by Hour of the Day

In [532...

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Sample DataFrame (Use your actual dataset)
# Ensure 'Date' and 'Time' are in proper datetime format
data = pd.DataFrame({
    'Date': pd.date_range(start='2020-01-01', periods=100, freq='H'),
    'Sales': [i % 24 + 1 for i in range(100)]
})


# Ensure 'Date' is in datetime format
data['Date'] = pd.to_datetime(data['Date'])

# Extract the hour from the Date column
df['Hour_of_day'] = data['Date'].dt.hour

# Group sales data by hour of the day
sales_by_time = df.groupby('Hour_of_day')['Sales'].sum().reset_index()

# Plotting
plt.figure(figsize=(12, 6))
sns.lineplot(x='Hour_of_day', y='Sales', data=sales_by_time, marker="o", color="b")
```

```
# Formatting the plot
plt.title("Sales Analysis by Hour of the Day", fontsize=16)
plt.xlabel("Hour of the Day (0-23)", fontsize=12)
plt.ylabel("Total Sales", fontsize=12)
plt.xticks(range(0, 24)) # Ensure all hours from 0 to 23 are displayed
plt.grid(True, linestyle="--", alpha=0.7)
plt.show()
```



In [533...

```
sales_by_time = df.groupby('Hour_of_day')['Sales'].sum().reset_index()
sales_by_time
```

Out[533]:

	Hour_of_day	Sales
0	0.0	0.761905
1	1.0	0.746032
2	2.0	0.650794
3	3.0	1.079365
4	4.0	0.571429
5	5.0	0.714286
6	6.0	0.587302
7	7.0	0.888889
8	8.0	0.587302
9	9.0	0.714286
10	10.0	0.634921
11	11.0	0.777778
12	12.0	1.634921
13	13.0	0.920635
14	14.0	1.222222
15	15.0	1.015873
16	16.0	1.365079
17	17.0	1.031746
18	18.0	0.984127
19	19.0	1.126984
20	20.0	0.968254
21	21.0	1.174603
22	22.0	1.365079
23	23.0	1.238095

In [534...]

```
df.head()
```

Out[534]:

	Date	index	Time	State	Group	Unit	Sales	Time_n	State_n	Group_n	Demog
0	2020-10-01	0	Morning	WA	Kids	0.095238	0.095238	2	6	0	
1	2020-10-01	1	Morning	WA	Men	0.095238	0.095238	2	6	1	V
2	2020-10-01	2	Morning	WA	Women	0.031746	0.031746	2	6	3	S
3	2020-10-01	3	Morning	WA	Seniors	0.206349	0.206349	2	6	2	
4	2020-10-01	4	Afternoon	WA	Kids	0.015873	0.015873	0	6	0	



Plotting Correlation matrix to identify the important features that affects the sales

```
In [543... df_numeric= df.drop(['Time', 'State', 'Group', 'Demographic', 'Region'],axis='columns')
df_numeric
```

```
Out[543]:
```

	Date	index	Unit	Sales	Time_n	State_n	Group_n	Hour_of_day
0	2020-10-01	0	0.095238	0.095238	2	6	0	0.0
1	2020-10-01	1	0.095238	0.095238	2	6	1	1.0
2	2020-10-01	2	0.031746	0.031746	2	6	3	2.0
3	2020-10-01	3	0.206349	0.206349	2	6	2	3.0
4	2020-10-01	4	0.015873	0.015873	0	6	0	4.0
...
7555	2020-12-30	7555	0.190476	0.190476	0	4	2	NaN
7556	2020-12-30	7556	0.206349	0.206349	1	4	0	NaN
7557	2020-12-30	7557	0.206349	0.206349	1	4	1	NaN
7558	2020-12-30	7558	0.142857	0.142857	1	4	3	NaN
7559	2020-12-30	7559	0.174603	0.174603	1	4	2	NaN

7560 rows × 8 columns

```
In [544... corr_matrix= df_numeric.corr()
corr_matrix
```

```
Out[544]:
```

	Date	index	Unit	Sales	Time_n	State_n	Group_n	Hour_of_day
Date	1.000000e+00	0.999885	0.100658	0.100658	2.593151e-15	2.481039e-15	-1.317944e-15	1.482944e-01
index	9.998850e-01	1.000000	0.104095	0.104095	-7.482612e-04	-3.534798e-03	4.098395e-03	0.131689
Unit	1.006582e-01	0.104095	1.000000	1.000000	1.004629e-03	-6.439166e-03	-1.105118e-03	0.309362
Sales	1.006582e-01	0.104095	1.000000	1.000000	1.004629e-03	-6.439166e-03	-1.105118e-03	0.309362
Time_n	2.593151e-15	-0.000748	0.001005	0.001005	1.000000e+00	-3.003558e-18	5.743231e-18	-2.860248e-01
State_n	2.481039e-15	-0.003535	-0.006439	-0.006439	-3.003558e-18	1.000000e+00	-3.599018e-18	-8.309249e-02
Group_n	-1.317944e-15	0.000410	-0.001105	-0.001105	5.743231e-18	-3.599018e-18	1.000000e+00	1.266305e-01
Hour_of_day	1.482944e-01	0.131689	0.309362	0.309362	-2.860248e-01	-8.309249e-02	1.266305e-01	1.000000e+00

```
In [545... plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", fmt=".2f", linewidths=0.5)
```

```
plt.title("Correlation Matrix Heatmap")
plt.show()
```



from the correlation matrix we can clearly see that only Unit and Hour of the day column/feature has a strong correlation with the Sales. But still we would confirm this with Hypothesis testing

```
In [546... df_numeric.head()
```

```
Out[546]:
```

	Date	index	Unit	Sales	Time_n	State_n	Group_n	Hour_of_day
0	2020-10-01	0	0.095238	0.095238	2	6	0	0.0
1	2020-10-01	1	0.095238	0.095238	2	6	1	1.0
2	2020-10-01	2	0.031746	0.031746	2	6	3	2.0
3	2020-10-01	3	0.206349	0.206349	2	6	2	3.0
4	2020-10-01	4	0.015873	0.015873	0	6	0	4.0

```
In [558... data = df_numeric.rename(columns=lambda x: x.strip()) # Removes Leading/trailing s
data.head()
```

Out[558]:

	Date	index	Unit	Sales	Time_n	State_n	Group_n	Hour_of_day
0	2020-10-01	0	0.095238	0.095238	2	6	0	0.0
1	2020-10-01	1	0.095238	0.095238	2	6	1	1.0
2	2020-10-01	2	0.031746	0.031746	2	6	3	2.0
3	2020-10-01	3	0.206349	0.206349	2	6	2	3.0
4	2020-10-01	4	0.015873	0.015873	0	6	0	4.0

Null Hypothesis (H0): "Units sold has nothing to do with sales revenue"

In [560]...

```
import scipy.stats as stats
import statsmodels.api as sm
import statsmodels.formula.api as smf

Unit_Sales_corr, Unit_Sales_pval = stats.spearmanr(data['Unit'], data['Sales'])

print(f"Spearman correlation between Unit and Sales: {Unit_Sales_corr:.2f}, p-value: {Unit_Sales_pval:.2f}")

# Regression analysis for age and shopping preferences
Unit_Sales_model = smf.ols('Sales ~ Unit', data=data).fit()

print("\nIn-store Purchases Regression Summary:\n", Unit_Sales_model.summary())
```

Spearman correlation between Unit and Sales: 1.00, p-value: 0.0000

In-store Purchases Regression Summary:

OLS Regression Results

```
=====
Dep. Variable:          Sales      R-squared:                1.000
Model:                  OLS      Adj. R-squared:            1.000
Method:                 Least Squares    F-statistic:          1.949e+33
Date:                  Sat, 01 Feb 2025    Prob (F-statistic):      0.00
Time:                  17:43:01    Log-Likelihood:         2.5725e+05
No. Observations:      7560    AIC:                   -5.145e+05
Df Residuals:          7558    BIC:                   -5.145e+05
Df Model:               1
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-5.345e-16	7.39e-18	-72.324	0.000	-5.49e-16	-5.2e-16
Unit	1.0000	2.26e-17	4.42e+16	0.000	1.000	1.000

```
=====
Omnibus:                 1514.899    Durbin-Watson:           0.054
Prob(Omnibus):            0.000    Jarque-Bera (JB):        2734.456
Skew:                    -1.271    Prob(JB):                0.00
Kurtosis:                 4.491    Cond. No.                 5.21
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Since the p-value is < 0.05 The result is highly statistically significant, meaning there is strong evidence to reject the null

hypothesis that "Units sold has nothing to do with sales revenue".

Model Construction

Input Variable (X): Unit and Hour of the day

```
In [ ]: X= df_numeric.drop(['Date', 'index', 'Sales', 'Time_n', 'State_n', 'Group_n'],axis= 'col')
X.head()
```

```
In [ ]: X['Hour_of_day'].isnull().sum()
```

```
In [ ]: X['Hour_of_day'].fillna(X['Hour_of_day'].mode()[0], inplace=True) # Fill with mode
```

Target Variable (Y) : Sales

```
In [ ]: Y= df['Sales']
Y.head()
```

Splitting the dataset into train and test parts where we are using 80% our data to train and 20% to test

```
In [ ]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test= train_test_split(X,Y,test_size=0.2, random_state=
```

Trying to plot a Linear regression model as we had previously seen that our data follows a linear pattern

```
In [ ]: from sklearn.linear_model import LinearRegression
reg= LinearRegression()
reg.fit(X_train, Y_train)
reg.score(X_test,Y_test)
```

The R2 score of our model resulted in a perfect fit.

```
In [ ]: Y_pred= reg.predict(X_test)
Y_pred
```

```
In [ ]: len(Y_pred)
```

```
In [ ]: len(Y)
```

```
In [ ]: Y = Y[:len(Y_pred)]
r2 = r2_score(Y_test, Y_pred)
```

```
In [ ]: from sklearn.metrics import mean_squared_error, r2_score
r2 = r2_score(Y_test, Y_pred)
print("R-squared Score:", r2)
```

Calculated the root mean square error as well to see how our model is performing

```
In [ ]: rmse = np.sqrt(mean_squared_error(Y_test, Y_pred))
rmse
```

As it can be seen that our Root Mean Squared Error (RMSE) = 3.42×10^{-16} is extremely close to zero, which aligns with our $R^2 = 1.0$. This means our model is predicting Y_{test} almost perfectly. However, this is highly unusual in real-world scenarios and often indicates overfitting. Hence just to confirm we are using K-fold cross validation and Ridge Regression

```
In [ ]: import numpy as np
from sklearn.model_selection import cross_val_score
reg= LinearRegression()
Score= cross_val_score(reg,X_test,Y_test,cv=3)
np.average(Score)
```

```
In [ ]: from sklearn.linear_model import Ridge
ridge_reg= Ridge(alpha=50, max_iter=100, tol=0.1)
ridge_reg.fit(X_train,Y_train)
ridge_reg.score(X_test,Y_test)
```

As it can be seen and confirmed from above that the model performs and fits properly even after perform cross validation and ridge regression. This indicates that our model is reliable.

```
In [ ]: from sklearn.linear_model import LinearRegression

# Instantiate the model
reg = LinearRegression()

# Fit the model on training data
reg.fit(X_train, Y_train)

# Now make predictions on the test data
Y_pred = reg.predict(X_test)

# You can now evaluate the performance using Y_pred
print("Predictions:", Y_pred)
```

```
In [540]: reg.predict([[0.015873,4.0]])
```

C:\Users\captr\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(

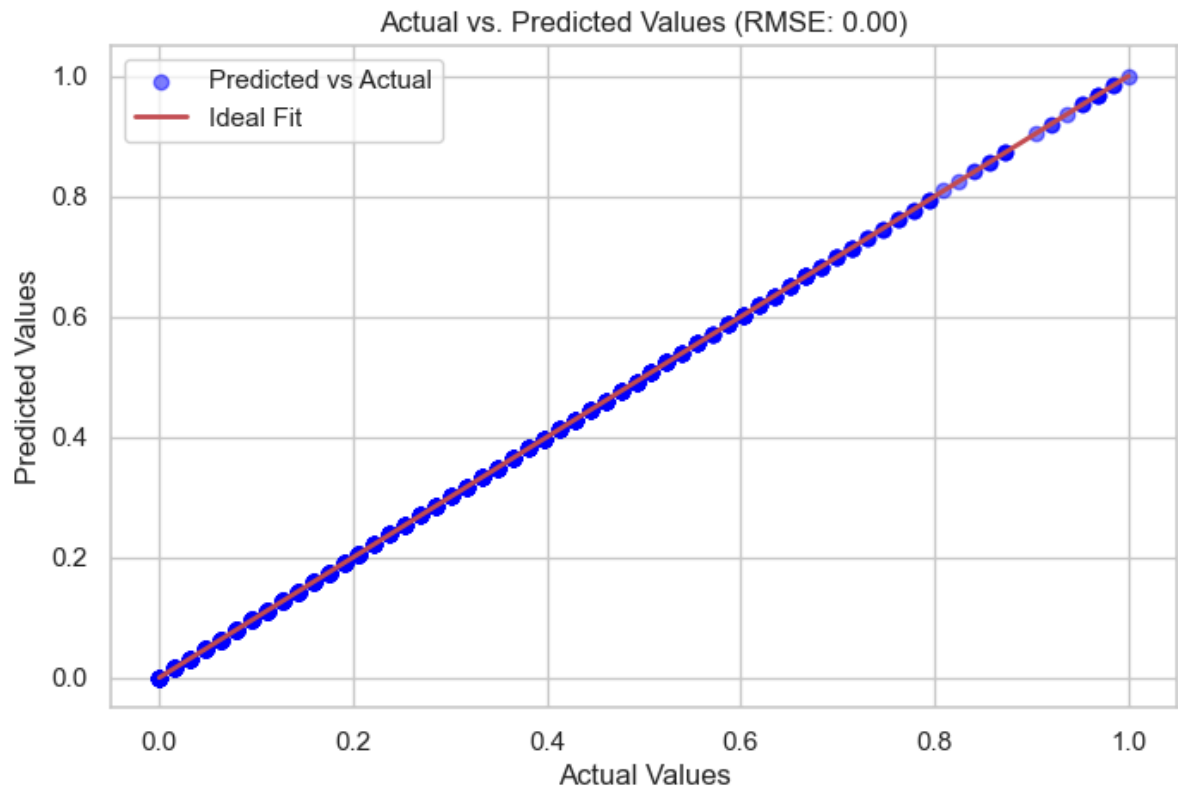
```
Out[540]: array([0.015873])
```

RMSE Plot

```
In [541]: plt.figure(figsize=(8, 5))
plt.scatter(Y_test, Y_pred, color="blue", alpha=0.5, label="Predicted vs Actual")
```



```
plt.plot([Y_test.min(), Y_test.max()], [Y_test.min(), Y_test.max()], 'r', lw=2, label="Actual Values")
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title(f"Actual vs. Predicted Values (RMSE: {rmse:.2f})")
plt.legend()
plt.show()
```




In [542...

```
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import mean_squared_error

# Compute RMSE
rmse = np.sqrt(mean_squared_error(Y_test, Y_pred))

# Plot RMSE
plt.figure(figsize=(6, 4))
plt.bar(["RMSE"], [rmse], color='blue')
plt.ylabel("Error")
plt.title("Root Mean Squared Error (RMSE)")
plt.ylim(0, rmse * 1.2) # Adding some padding for better visualization
plt.show()
```



```
In [ ]:
```